

Texture Quilting (Due Saturday 2/23/2019)

In this assignment, you will develop code to stitch together image patches sampled from an input texture in order to synthesize new texture images. You can download the test image used to generate the example above from assignment folder Canvas.

You should start by reading through the whole assignment, looking at the provided code in detail to make sure you understand what it does. The main function ***quilt_demo*** appears at the end. You will need to write several subroutines in order for it to function properly.

Name: Dikai Fang

SID: 29991751

1. Shortest Path [25 pts]

Write a function ***shortest_path*** that takes an 2D array of ***costs***, of shape HxW, as input and finds the *shortest vertical path* from top to bottom through the array. A vertical path is specified by a single horizontal location for each row of the H rows. Locations in successive rows should not differ by more than 1 so that at each step the path either goes straight or moves at most one pixel to the left or right. The cost is the sum of the costs of each entry the path traverses. Your function should return an length H vector that contains the index of the path location (values in the range 0..W-1) for each of the H rows.

You should solve the problem by implementing the dynamic programming algorithm described in class. You will have a for-loop over the rows of the "cost-to-go" array (M in the slides), computing the cost of the shortest path up to that row using the recursive formula that depends on the costs-to-go for the previous row. Once you have get to the last row, you can then find the smallest total cost. To find the path which actually has this smallest cost, you will need to do backtracking. The easiest way to do this is to also store the index of whichever minimum was selected at each location. These indices will also be an HxW array. You can then backtrack through these indices, reading out the path.

Finally, you should create at least three test cases by hand where you know the shortest path and see that the code gives the correct answer.

```
In [180]: 1 #modules used in your code
          2 import numpy as np
          3 import matplotlib.pyplot as plt
```

```
In [181]: 1 def shortest_path(costs):
          2     """
          3     This function takes an array of costs and finds a shortest path from the
          4     top to the bottom of the array which minimizes the total costs along the
          5     path. The path should not shift more than 1 location left or right between
          6     successive rows.
          7
          8     In other words, the returned path is one that minimizes the total cost:
          9
          10         total_cost = costs[0,path[0]] + costs[1,path[1]] + costs[2,path[2]] + ...
          11
          12     subject to the constraint that:
```

```

13
14     abs(path[i]-path[i+1])<=1
15
16 Parameters
17 -----
18 costs : 2D float array of shape HxW
19         An array of cost values
20
21 Returns
22 -----
23 path : 1D array of length H
24         indices of a vertical path. path[i] contains the column index of
25         the path for each row i.
26     """
27
28     pathMatrix = np.zeros(costs.shape, dtype = np.float64)
29     costMatrix = np.copy(pathMatrix)
30
31     h = costs.shape[0]
32     w = costs.shape[1]
33
34     for i in range(h-1,-1,-1):
35         if (i == h-1):
36             costMatrix[i] = costs[i]
37             pathMatrix[i] = np.array(range(w))
38         else:
39             inf = np.array(np.inf)
40             mid = np.copy(costMatrix[i+1])
41             left = np.concatenate((inf, mid[:w-1:]), axis = None)
42             right = np.concatenate((mid[1::], inf), axis = None)
43
44             threeCompare = np.stack((left, mid, right))
45             minStep = np.amin(threeCompare, axis=0)
46             costMatrix[i] = costs[i] + minStep
47
48             minIndices = np.argmin(threeCompare, axis = 0) - 1
49             pathMatrix[i] = np.array(range(w)) + minIndices
50
51     path = list()
52     path.append(np.argmin(costMatrix[0]))
53

```

```
53
54     for i in range(h-1):
55         p = path[i]
56         nextP = pathMatrix[i][int(p)]
57         path.append(int(nextP))
58
59     return path
```

In [182]:

```
1 #
2 # your test code goes here. come up with at least 3 test cases
3 #
4
5 costs1 = np.array([[1,2,2,2], [2,1,2,2], [2,2,1,2], [2,2,2,1]])
6 path1 = shortest_path(costs1)
7 print(costs1)
8 print(path1)
9
10 costs2 = np.array([[3,3,1,2,2],[3,3,1,4,5],[5,5,5,1,4],[6,6,6,1,4]])
11 path2 = shortest_path(costs2)
12
13 print(costs2)
14 print(path2)
15
16 costs3 = np.array([[9, 3, 12, 19, 3, 23, 0, 900], [30, 12, 34, 78, 2, 0, 1, 10], [20, 789, 3, 1
17 path3 = shortest_path(costs3)
18
19 print(costs3)
20 print(path3)
```

```
[[1 2 2 2]
 [2 1 2 2]
 [2 2 1 2]
 [2 2 2 1]]
[0, 1, 2, 3]
[[3 3 1 2 2]
 [3 3 1 4 5]
 [5 5 5 1 4]
 [6 6 6 1 4]]
[2, 2, 3, 3]
[[ 9  3 12 19  3 23  0 900]
 [30 12 34 78  2  0  1 10]
 [20 789  3 12 45 78 39 66]
 [12 23 98 78 34 49 79 12]]
[1, 1, 2, 1]
```

2. Image Stitching: [25 pts]

Write a function ***stitch*** that takes two gray-scale images, ***left_image*** and ***right_image*** and a specified ***overlap*** and returns a new output image by stitching them together along a seam where the two images have very similar brightness values. If the input images are of widths ***w1*** and ***w2*** then your stitched result image returned by the function should be of width ***w1+w2-overlap*** and have the same height as the two input images.

You will want to first extract the overlapping strips from the two input images and then compute a cost array given by the absolute value of their difference. You can then use your ***shortest_path*** function to find the seam along which to stitch the images where they differ the least in brightness. Finally you need to generate the output image by using pixels from the left image on the left side of the seam and from the right image on the right side of the seam. You may find it easiest to code this by first turning the path into an alpha mask for each image and then using the standard equation for compositing.

In [183]:

```
1 def stitch(left_image, right_image, overlap):
2     """
3     This function takes a pair of images with a specified overlap and stitches them
4     together by finding a minimal cost seam in the overlap region.
5
6     Parameters
7     -----
8     left_image : 2D float array of shape HxW1
9         Left image to stitch
10
11     right_image : 2D float array of shape HxW2
12         Right image to stitch
13
14     overlap : int
15         Width of the overlap zone between left and right image
16
17     Returns
18     -----
19     stitched : 2D float array of shape Hx(W1+W2-overlap)
20         The resulting stitched image
21     """
22
23     # inputs should be the same height
24     assert(left_image.shape[0]==right_image.shape[0])
```

```

25
26     # your code here
27     h = left_image.shape[0]
28     w1 = left_image.shape[1]
29     w2 = right_image.shape[1]
30
31     olArea = np.abs((left_image[:, w1-overlap:]-right_image[:, :overlap]))
32
33     path = shortest_path(olArea)
34
35     stitched = np.zeros((h, w1+w2-overlap))
36     for i in range(h):
37         p = path[i]
38         l = left_image[i][:w1-overlap+p:]
39         m = left_image[i][w1-overlap+p]
40         r = right_image[i][p+1:]
41         new = np.concatenate((l,m,r), axis = None)
42         stitched[i] = new
43
44     assert(stitched.shape[0]==left_image.shape[0])
45     assert(stitched.shape[1]==(left_image.shape[1]+right_image.shape[1]-overlap))
46     return stitched
47
48

```

In []:

1

3. Texture Quilting: [25 pts]

Write a function ***synth_quilt*** that takes as input an array indicating the set of texture tiles to use, an array containing the set of available texture tiles, the ***tilesize*** and ***overlap*** parameters and synthesizes the output texture by stitching together the tiles. ***synth_quilt*** should utilize your stitch function repeatedly. First, for each horizontal row of tiles, construct the stitched row by repeatedly stitching the next tile in the row on to the right side of your row image. Once you have row images for all the rows, you can stitch them together to get the final image. Since your stitch function only works for vertical seams, you will want to transpose the rows, stitch them together, and then transpose the result. You may find it useful to look at the provided code below which simply puts down the tiles with the specified overlap but doesn't do stitching. Your quilting function will return a similar result but with much smoother transitions between the tiles.

In [184]:

```
1  def synth_quilt(tile_map,tiledb,tilesize,overlap):
2
3      """
4      This function takes as input an array indicating the set of texture tiles
5      to use at each location, an array containing the database of available texture
6      tiles, tilesize and overlap parameters, and synthesizes the output texture by
7      stitching together the tiles
8
9
10     Parameters
11     -----
12     tile_map : 2D array of int
13         Array storing the indices of which tiles to paste down at each output location
14
15     tiledb : 2D array of size ntiles x npixels
16         Collection of sample tiles to select from
17
18     tilesize : (int,int)
19         Size of a tile in pixels
20
21     overlap : int
22         Amount of overlap between tiles
23
24     Returns
25     -----
26     output : 2D float array
```



```

27         The resulting synthesized texture of size
28         """
29
30         # determine output size based on overlap and tile size
31         outh = (tilesize[0]-overlap)*tile_map.shape[0] + overlap
32         outw = (tilesize[1]-overlap)*tile_map.shape[1] + overlap
33         output = np.zeros((outh,outw))
34
35         # The code below is a dummy implementation that pastes down each
36         # tile in the correct position in the output image. You need to
37         # replace this with your own version that stitches each row and then
38         # stitches together the columns
39
40         for i in range(tile_map.shape[0]):
41             tmp = np.zeros((outh, outw))
42             for j in range(tile_map.shape[1]):
43                 icoord = i*(tilesize[0]-overlap)
44                 jcoord = j*(tilesize[1]-overlap)
45                 tile_vec = tiledb[tile_map[i,j],:];
46                 tile_image = np.reshape(tile_vec,tilesize)
47
48
49                 if (j>0):
50                     p_jcoord = (j-1)*(tilesize[1]-overlap)
51                     left = tmp[icoord:(icoord+tilesize[0]), p_jcoord:p_jcoord+tilesize[1]]
52                     stitched_r = stitch(left, tile_image, overlap)
53                     tmp[icoord:(icoord+tilesize[0]), p_jcoord:(jcoord+tilesize[1])] = stitched_r
54                 else:
55                     tmp[icoord:(icoord+tilesize[0]),jcoord:(jcoord+tilesize[1])] = tile_image
56
57                 if (i>0):
58                     icoord = i*(tilesize[0]-overlap)
59                     p_icoord = (i-1)*(tilesize[0]-overlap)
60                     left = np.transpose(output[p_icoord:p_icoord+tilesize[0], :])
61                     right = np.transpose(tmp[icoord:icoord+tilesize[0], :])
62                     stitched_rows = stitch(left, right, overlap)
63                     tmp[p_icoord:(icoord+tilesize[0]), :] = np.transpose(stitched_rows)
64
65             output[icoord:(icoord+tilesize[0]), :] = tmp[icoord:(icoord+tilesize[0]), :]
66
67         return output

```

```
57  
68 return output
```

In []:

1

4. Texture Synthesis Demo [25pts]

The function provided below **quilt_demo** puts together the pieces. It takes a sample texture image and a specified output size and uses the functions you've implemented previously to synthesize a new texture sample.

You should write some additional code in the cells that follow to in order demonstrate the final result and experiment with the algorithm parameters in order to produce a compelling visual result and write explanations of what you discovered.

Test your code on the provided image *rock_wall.jpg*. There are three parameters of the algorithm. The *tilesize*, *overlap* and *K*. In the provided `*texture_demo**` code below, these have been set at some default values. Include in your demo below images of three example texture outputs when you: (1) increase the tile size, (2) decrease the overlap, and (3) decrease the value for *K*. For each result explain how it differs from the default setting of the parameters and why.

Test your code on two other texture source images of your choice. You can use images from the web or take a picture of a texture yourself. You may need to resize or crop your input image to make sure that the **tiledb** is not overly large. You will also likely need to modify the **tilesize** and **overlap** parameters depending on your choice of texture. Once you have found good settings for these parameters, synthesize a nice output texture. Make sure you display both the image of the input sample and the output synthesis for your two other example textures in your submitted pdf.

In [185]:

```
1  #skimage is only needed for sample tiles code provided below  
2  #you should not use it in your own code  
3  import skimage as ski  
4  
5  def sample_tiles(image,tilesize,randomize=True):  
6      """  
7          This function generates a library of tiles of a specified size from a given source image  
8  
9          Parameters  
10         -----  
11         image : float array of shape HxW  
12             Input image  
13
```

```

14     tileSize : (int,int)
15         Dimensions of the tiles in pixels
16
17
18     Returns
19     -----
20     tiles : float array of shape  numtiles x npixels
21         The library of tiles stored as vectors where npixels is the
22         product of the tile height and width
23     """
24
25     tiles = ski.util.view_as_windows(image,tileSize)
26     ntiles = tiles.shape[0]*tiles.shape[1]
27     npix = tiles.shape[2]*tiles.shape[3]
28     assert(npix==tileSize[0]*tileSize[1])
29
30     print("library has ntiles = ",ntiles,"each with npix = ",npix)
31
32     tiles = tiles.reshape((ntiles,npix))
33
34     # randomize tile order
35     if randomize:
36         tiles = tiles[np.random.permutation(ntiles),:]
37
38     return tiles
39
40
41 def topkmatch(tilestrip,dbstrips,k):
42     """
43     This function finds the top k candidate matches in dbstrips that
44     are most similar to the provided tile strip.
45
46     Parameters
47     -----
48     tilestrip : 1D float array of length npixels
49         Grayscale values of the query strip
50
51     dbstrips : 2D float array of size npixels x numtiles
52         Array containing brightness values of numtiles strips in the database
53         to match to the npixels brightness values in tilestrip
54     """

```

```

54
55     k : int
56         Number of top candidate matches to sample from
57
58     Returns
59     -----
60     matches : list of ints of length k
61         The indices of the k top matching tiles
62     """
63     assert(k>0)
64     assert(dbstrips.shape[0]>k)
65     error = (dbstrips-tilestrip)
66     ssd = np.sum(error*error,axis=1)
67     ind = np.argsort(ssd)
68     matches = ind[0:k]
69     return matches
70
71
72 def quilt_demo(sample_image, ntilesout=(10,20), tileSize=(30,30), overlap=5, k=5):
73     """
74     This function takes an image and quilting parameters and synthesizes a
75     new texture image by stitching together sampled tiles from the source image.
76
77
78     Parameters
79     -----
80     sample_image : 2D float array
81         Grayscale image containing sample texture
82
83     ntilesout : list of int
84         Dimensions of output in tiles, e.g. (3,4)
85
86     tileSize : int
87         Size of the square tile in pixels
88
89     overlap : int
90         Amount of overlap between tiles
91
92     k : int
93         Number of top candidate matches to sample from
94

```

```

95     Returns
96     -----
97     img : list of int of length K
98           The resulting synthesized texture of size
99     """
100
101     # generate database of tiles from sample
102     tiledb = sample_tiles(sample_image,tilesize)
103     # number of tiles in the database
104     nsampletiles = tiledb.shape[0]
105
106     if (nsampletiles<k):
107         print("Error: tile database is not big enough!")
108
109     # generate indices of the different tile strips
110     i,j = np.mgrid[0:tilesize[0],0:tilesize[1]]
111     top_ind = np.ravel_multi_index(np.where(i<overlap),tilesize)
112     bottom_ind = np.ravel_multi_index(np.where(i>=tilesize[0]-overlap),tilesize)
113     left_ind = np.ravel_multi_index(np.where(j<overlap),tilesize)
114     right_ind = np.ravel_multi_index(np.where(j>=tilesize[1]-overlap),tilesize)
115
116     # initialize an array to store which tile will be placed
117     # in each location in the output image
118     tile_map = np.zeros(ntilesout,'int')
119
120     # print('row:')
121     for i in range(ntilesout[0]):
122         # print(i)
123         for j in range(ntilesout[1]):
124
125             if (i==0)&(j==0):                # first row first tile
126                 matches = np.zeros(k) #range(nsampletiles)
127
128             elif (i==0):                    # first row (but not first tile)
129                 left_tile = tile_map[i,j-1]
130                 tilestrip = tiledb[left_tile,right_ind]
131                 dbstrips = tiledb[:,left_ind]
132                 matches = topkmatch(tilestrip,dbstrips,k)
133
134             elif (j==0):                    # first column (but not first row)
135                 above_tile = tile_map[i-1,j]

```

```

136         tilestrip = tiledb[above_tile,bottom_ind]
137         dbstrips = tiledb[:,top_ind]
138         matches = topkmatch(tilestrip,dbstrips,k)
139
140     else:                                     # neighbors above and to the left
141         left_tile = tile_map[i,j-1]
142         tilestrip_1 = tiledb[left_tile,right_ind]
143         dbstrips_1 = tiledb[:,left_ind]
144         above_tile = tile_map[i-1,j]
145         tilestrip_2 = tiledb[above_tile,bottom_ind]
146         dbstrips_2 = tiledb[:,top_ind]
147         # concatenate the two strips
148         tilestrip = np.concatenate((tilestrip_1,tilestrip_2))
149         dbstrips = np.concatenate((dbstrips_1,dbstrips_2),axis=1)
150         matches = topkmatch(tilestrip,dbstrips,k)
151
152         #choose one of the k matches at random
153         tile_map[i,j] = matches[np.random.randint(0,k)]
154
155     output = synth_quilt(tile_map,tiledb,tilesize,overlap)
156
157     return output
158

```

In [186]:

```

1  # load in rock_wall.jpg
2  # run and display results for quilt_demo with
3  #
4  # (0) default parameters
5  # (1) increased tile size
6  # (2) decrease the overlap
7  # (3) increase the value for K.
8
9  I = plt.imread("rock_wall.jpg")
10 grayI = np.zeros((I.shape[0], I.shape[1]), dtype = np.float32)
11 grayI[:, :, :] = (0.299 * I[:, :, 0] + 0.587 * I[:, :, 1] + 0.114 * I[:, :, 2])
12 I = grayI
13
14 output = quilt_demo(I)
15 output2 = quilt_demo(I, tilesize = (60,60))
16 output3 = quilt_demo(I, overlap = 2)

```

```




17 output4 = quilt_demo(I, k= 10)
18
19 fig = plt.figure(figsize=(20,19))
20 fig.add_subplot(1,2,1).imshow(output,cmap=plt.cm.gray)
21 fig.add_subplot(1,2,2).imshow(output2,cmap=plt.cm.gray)
22
23 fig2 = plt.figure(figsize=(20,19))
24 fig2.add_subplot(1,2,1).imshow(output3,cmap=plt.cm.gray)
25 fig2.add_subplot(1,2,2).imshow(output4,cmap=plt.cm.gray)
26
27 plt.show()
28


```

```

library has ntiles = 29241 each with npix = 900
library has ntiles = 19881 each with npix = 3600
library has ntiles = 29241 each with npix = 900
library has ntiles = 29241 each with npix = 900

```





For each result shown, explain here how it differs visually from the default setting of the parameters and explain why:

- . increase tile size will make the tile contain more content: rocks. Therefore the result image's content is denser
- . decrease overlap reduce the possibility of path and limit the minimal cost. Therefore the border between tiles is more obvious
- . increase value of k means to give more possibilities, which also means to increase the possibilities of errors. Therefore the result is not good as the default
- .

In [187]:

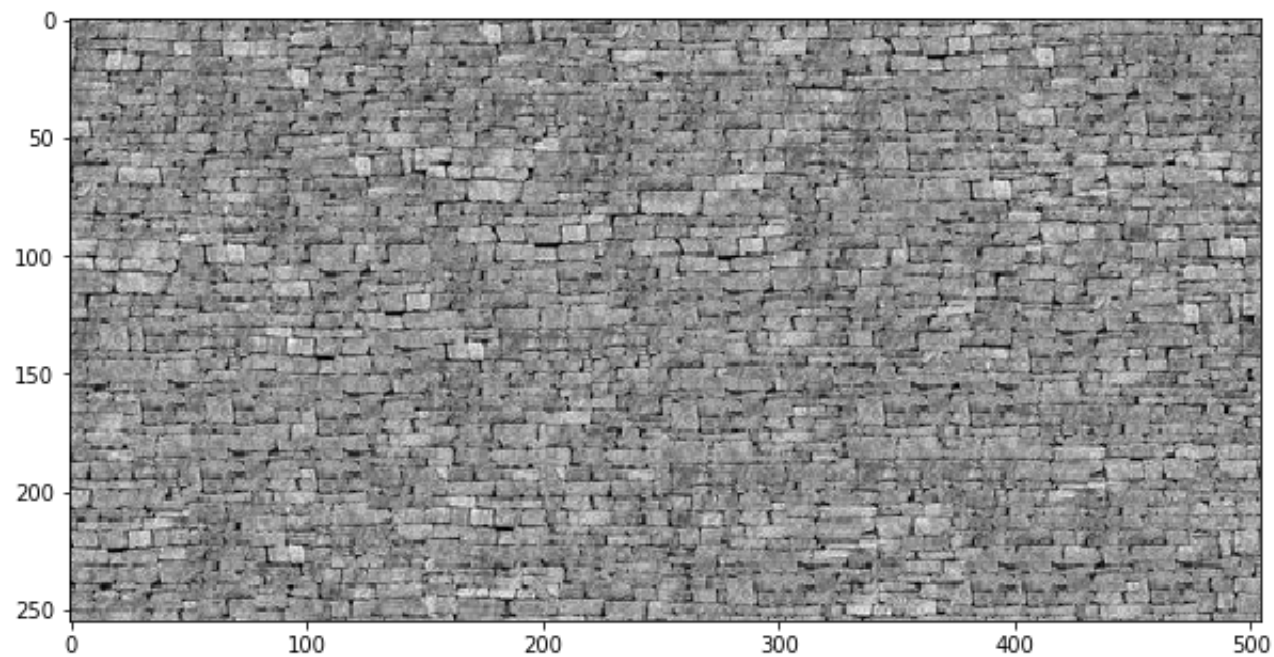
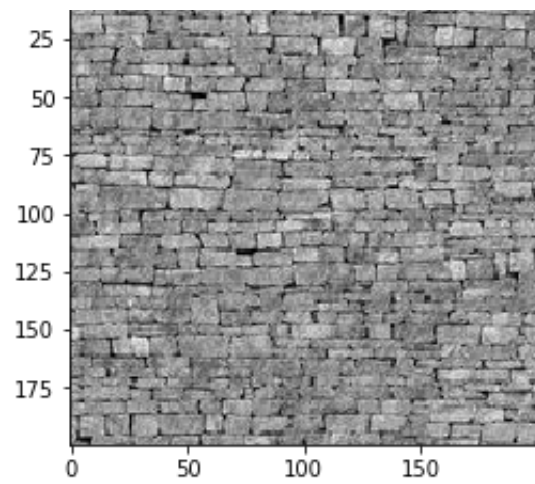
```

1  #
2  # load in yourimage1.jpg
3  #
4  # call quilt_demo, experiment with parameters if needed to get a good result
5  #
6  # display your source image and the resulting synthesized texture
7  #
8
9  I = plt.imread("wall.jpg")
10 grayI = np.zeros((I.shape[0], I.shape[1]), dtype = np.float32)
11 grayI[:, :, :] = (0.299 * I[:, :, 0] + 0.587 * I[:, :, 1] + 0.114 * I[:, :, 2])
12
13 oriFig = plt.figure()
14 oriFig.add_subplot(1,1,1).imshow(grayI,cmap=plt.cm.gray)
15
16 synthFig = plt.figure(figsize=(10,9))
17 output = quilt_demo(grayI)
18 synthFig.add_subplot(1,1,1).imshow(output,cmap=plt.cm.gray)
19
20 plt.show()
21
22

```

library has ntiles = 29241 each with npix = 900





In [188]:



```
1 #  
2 # load in yourimage2.jpg  
3 #  
4 # call quilt_demo, experiment with parameters if needed to get a good result  
5 #  
6 # display your source image and the resulting synthesized texture  
7 #
```


```

8
9 I = plt.imread("water.jpg")
10 grayI = np.zeros((I.shape[0], I.shape[1]), dtype = np.float32)
11 grayI[:, :, :] = (0.299 * I[:, :, 0] + 0.587 * I[:, :, 1] + 0.114 * I[:, :, 2])
12
13 oriFig = plt.figure()
14 oriFig.add_subplot(1,1,1).imshow(grayI,cmap=plt.cm.gray)
15
16 synthFig = plt.figure(figsize=(10,9))
17 output = quilt_demo(grayI, overlap = 4, k=3)
18 synthFig.add_subplot(1,1,1).imshow(output,cmap=plt.cm.gray)
19
20 plt.show()
21

```

library has ntiles = 29241 each with npix = 900





In []:

1

In []:

1