

Face Morphing and Swapping (due Sunday 3/17/2019)

In this assignment, you will develop a function to warp from one face to another using the piecewise affine warping technique described in class and use it to perform morphing and face-swapping.

Name: Dikai Fang

SID: 29991751

```
In [254]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import pickle
4
5 #part 2
6 from matplotlib.path import Path
7 from scipy.spatial import Delaunay
8 from a5utils import bilinear_interpolate
9
10 #part 2 demo for displaying animations in notebook
11 from IPython.display import HTML
12 from a5utils import display_movie
13
14 #part 4 blending
15 from scipy.ndimage import gaussian_filter
```

1. Transforming Triangles [30 pts]

Write a function ***get_transform*** which takes the corner coordinates of two triangles and computes an affine transformation (represented as a 3x3 matrix) that maps the vertices of a given source triangle to the specified target position. We will use this to map pixels inside each triangle of our mesh. For convenience, you should implement a function ***apply_transform*** that takes a transformation (3x3 matrix) and a set of points, and transforms the points.

```
In [255]: 1 def get_transform(pts_source,pts_target):
2     """
3     This function takes the coordinates of 3 points (corners of a triangle)
4     and a target position and estimates the affine transformation needed
5     to map the source to the target location.
6
7
8     Parameters
9     -----
10    pts_source : 2D float array of shape 2x3
11        Source point coordinates
12    pts_target : 2D float array of shape 2x3
```

```

13         Target point coordinates
14
15     Returns
16     -----
17     T : 2D float array of shape 3x3
18         the affine transformation
19     """
20
21     assert(pts_source.shape==(2,3))
22     assert(pts_source.shape==(2,3))
23
24     # your code goes here (see lecture #16)
25     # T = np.zeros((2,3), dtype = float)
26
27     oneRow = np.ones((1,3))
28     _pts_source = np.concatenate((pts_source, oneRow), axis = 0)
29     _pts_target = np.concatenate((pts_target, oneRow), axis = 0)
30
31     try:
32         inv_source = np.linalg.inv(_pts_source)
33     except np.linalg.LinAlgError:
34         print("matrix is not invertible")
35         return
36
37     T = np.matmul(_pts_target, inv_source)
38
39     return T
40
41
42 def apply_transform(T,pts):
43     """
44     This function takes the coordinates of a set of points and
45     a 3x3 transformation matrix T and returns the transformed
46     coordinates
47
48
49     Parameters
50     -----
51     T : 2D float array of shape 3x3
52         Transformation matrix
53     pts : 2D float array of shape 2xN

```

```

54         Set of points to transform
55
56     Returns
57     -----
58     pts_warped : 2D float array of shape 2xN
59         Transformed points
60     """
61
62     assert(T.shape==(3,3))
63     assert(pts.shape[0]==2)
64
65     # convert to homogenous coordinates, multiply by T, convert back
66     r = np.zeros((1, pts.shape[1]))
67     r[-1] = 1.
68     pts_h = np.concatenate((pts, r), axis = 0)
69
70     pts_warped = np.matmul(T, pts_h)
71     pts_warped = np.delete(pts_warped, (2), axis = 0)
72     # print(pts_warped)
73
74     return pts_warped

```

In [256]:

```
1  #
2  # Write some test cases for your affine_transform function
3  #
4
5  # check that using the same source and target should yield identity matrix
6  print("Same source and target: ")
7  src = np.array([[1,2,3],[2,3,9]])
8  targ = src
9  print(get_transform(src,targ))
10
11 # check that if targ is just a translated version of src, then the translation
12 # appears in the expected locations in the transformation matrix
13 print("Translated: ")
14 src = np.array([[1,2,3], [2,3,9]])
15 targ = np.array([[2,3,4], [3,4,10]])
16 print(get_transform(src,targ))
17
18 # random tests... check that for two random
19 # triangles the estimated transformation correctly
20 # maps one to the other
21 print("Random: ")
22 for i in range(5):
23     src = np.random.random((2,3))
24     targ = np.random.random((2,3))
25     T = get_transform(src,targ)
26     targ1 = apply_transform(T,src)
27     assert(np.sum(np.abs(targ-targ1))<1e-12)
28
```

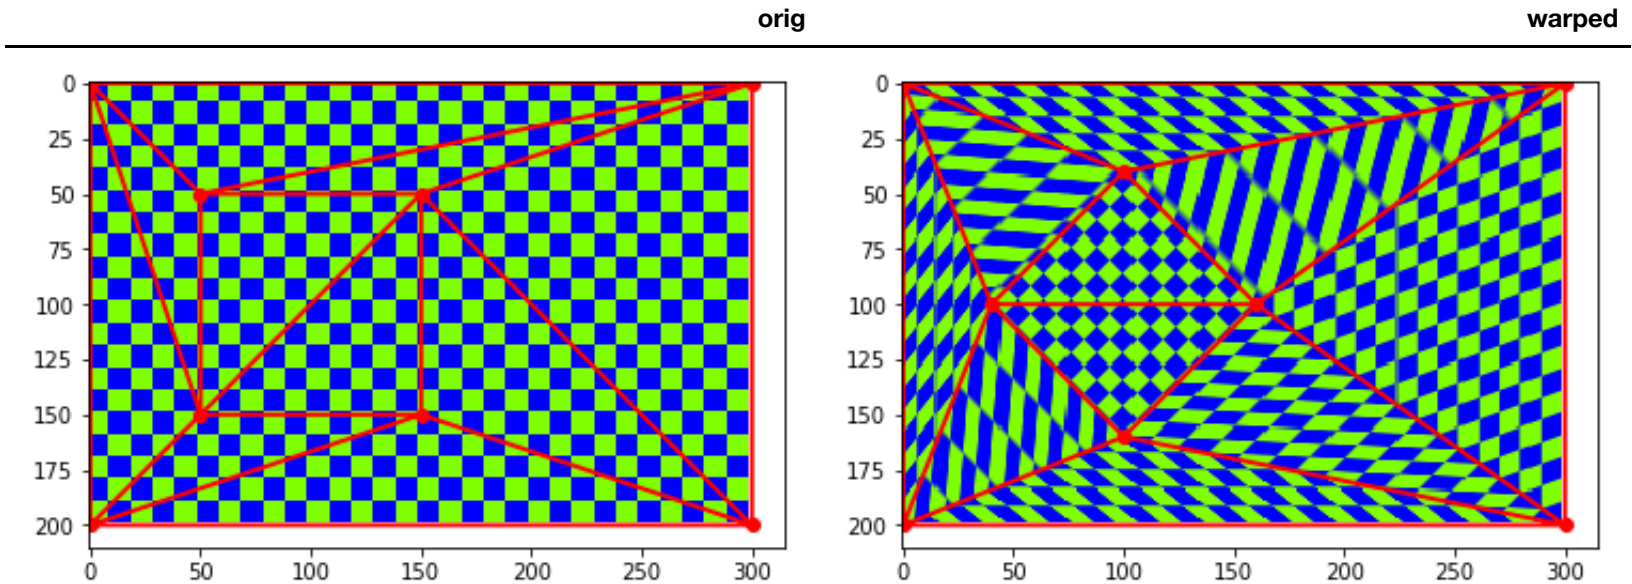
Same source and target:

```
[[ 1.00000000e+00  5.55111512e-17  5.55111512e-17]
 [-2.77555756e-16  1.00000000e+00  1.66533454e-16]
 [ 1.66533454e-16 -2.77555756e-17  1.00000000e+00]]
```

Translated:

```
[[ 1.00000000e+00  0.00000000e+00  1.00000000e+00]
 [ 3.33066907e-16  1.00000000e+00  1.00000000e+00]
 [ 1.66533454e-16 -2.77555756e-17  1.00000000e+00]]
```

Random:



2. Piecewise Affine Warping [40 pts]

Write a function called **warp** that performs piecewise affine warping of the image. Your function should take a source image, a set of triangulated points in the source image and a set of target locations for those points. We will accomplish this using *backwards warping* in the following steps:

1. For each pixel in the warped output image, you first need to determine which triangle it falls inside of. For this we can use **matplotlib.path.Path.contains_points** which checks whether a point falls inside a specified polygon. Your code should build an array **tindex** which is the same size as the input image where **tindex[i,j]=t** if pixel **[i,j]** falls inside triangle **t**. Pixels which are not in any triangle should have a **tindex** value of -1.
2. For each triangle, use your **get_transform** function from part 1 to compute the affine transformation which maps the pixels in the output image back to the source image (i.e., mapping **pts_target** to **pts_source** for the triangle). Apply the estimated transform to the coordinates of all the pixels in the output triangle to determine their locations in the input image.
3. Use bilinear interpolation to determine the colors of the output pixels. The provided code **a5utils.py** contains a function **bilinear_interpolate** that implements the interpolation. To handle color images, you will need to call **bilinear_interpolate** three times for the R, G and B color channels separately.

```

In [257]: 1 def warp(image,pts_source,pts_target,tri):
2
3         """
4         This function takes a color image, a triangulated set of keypoints
5         over the image, and a set of target locations for those points.
6         The function performs piecewise affine warping by warping the
7         contents of each triangle to the desired target location and
8         returns the resulting warped image.
9
10        Parameters
11        -----
12        image : 3D float array of shape HxWx3
13               An array containing a color image
14
15        pts_src: 2D float array of shape 2xN
16               Coordinates of N points in the image
17
18        pts_target: 2D float array of shape 2xN
19               Coordinates of the N points after warping
20
21        tri: 2D int array of shape Ntri x 3
22             The indices of the pts belonging to each of the Ntri triangles
23
24        Returns
25        -----
26        warped_image : 3D float array of shape HxWx3
27                       resulting warped image
28
29        tindex : 2D int array of shape HxW
30                array with values in 0...Ntri-1 indicating which triangle
31                each pixel was contained in (or -1 if the pixel is not in any triangle)
32        """
33
34        assert(image.shape[2]==3) #this function only works for color images
35        assert(tri.shape[1]==3)    #each triangle has 3 vertices
36        assert(pts_source.shape==pts_target.shape)
37        assert(np.max(image)<=1)   #image should be float with RGB values in 0..1
38
39        ntri = tri.shape[0]
40        (h,w,d) = image.shape
41

```

```

42     # for each pixel in the target image, figure out which triangle
43     # it fall in side of so we know which transformation to use for
44     # those pixels.
45     #
46     # tindex[i,j] should contain a value in 0..ntri-1 indicating which
47     # triangle contains pixel (i,j). set tindex[i,j]=-1 if (i,j) doesn't
48     # fall inside any triangle
49     tindex = -1*np.ones((h,w))
50     xx,yy = np.mgrid[0:h,0:w]
51     pcoords = np.stack((yy.flatten(),xx.flatten()),axis=1)
52     for t in range(ntri):
53         # corners = ... #Vertices of triangle t. Path expects a Kx2 array of vertices as inp
54         corners = np.array([[pts_target[0][tri[t][0]], pts_target[1][tri[t][0]]],
55                             [pts_target[0][tri[t][1]], pts_target[1][tri[t][1]]],
56                             [pts_target[0][tri[t][2]], pts_target[1][tri[t][2]]]])
57
58         path = Path(corners)
59         mask = path.contains_points(pcoords)
60         mask = mask.reshape(h,w)
61         #set tindex[i,j]=t any where that mask[i,j]=True
62         tindex[np.where(mask == True)] = t
63
64     # compute the affine transform associated with each triangle that
65     # maps a given target triangle back to the source coordinates
66
67     Xsource = np.zeros((2,h*w)) #source coordinate for each output pixel
68     tindex_flat = tindex.flatten() #flattened version of tindex as an h*w length vector
69
70     for t in range(ntri):
71         #coordinates of target/output vertices of triangle t
72         # targ = ...
73         targ = np.array([pts_target[0,tri[t]],pts_target[1,tri[t]]])
74
75
76         #coordinates of source/input vertices of triangle t
77         # psrc = ...
78         psrc = np.array([pts_source[0,tri[t]],pts_source[1,tri[t]]])
79
80         #compute transform from ptarg -> psrc
81         # T = ...
82         T = get_transform(targ, psrc)

```



```

82         get_transform(img, psrc,
83
84     #extract coordinates of all the pixels where tindex==t
85     #
86     pcoords_t = ...
87     indices = np.where(tindex==t)
88     pcoords_t = np.vstack((indices[1], indices[0]))
89     #
90     pcoords_t = pcoords[np.where(tindex_flat == t)]
91
92     #store the transformed coordinates at the corresponding locations in Xsource
93     #
94     print(T.shape)
95     #
96     print(T)
97     #
98     print(pcoords_t.shape)
99     #
100    print(pcoords_t)
101    Xsource[:,tindex_flat==t] = apply_transform(T,pcoords_t)
102    #
103    Xsource[:,tindex==t] -> Xsource[:,np.where(tindex==t)]
104
105    # now use interpolation to figure out the color values at locations Xsource
106    warped_image = np.zeros(image.shape)
107    warped_image[:, :, 0] = bilinear_interpolate(image[:, :, 0], Xsource[0, :], Xsource[1, :]).reshape(Xsource[0, :].shape)
108    warped_image[:, :, 1] = bilinear_interpolate(image[:, :, 1], Xsource[0, :], Xsource[1, :]).reshape(Xsource[0, :].shape)
109    warped_image[:, :, 2] = bilinear_interpolate(image[:, :, 2], Xsource[0, :], Xsource[1, :]).reshape(Xsource[0, :].shape)
110
111    # clip RGB values outside the range [0,1] to avoid warning messages
112    # when displaying warped image later on
113    warped_image = np.clip(warped_image, 0., 1.)
114
115    return (warped_image, tindex)

```

In [258]:

```

1  #
2  # Test your warp function
3  #
4
5  #make a color checkerboard image
6  (xx,yy) = np.mgrid[1:200,1:300]
7  G = np.mod(np.floor(xx/10)+np.floor(yy/10),2)
8  B = np.mod(np.floor(xx/10)+np.floor(yy/10)+1,2)
9  image = np.stack((0.5*G,G,B),axis=2)
10
11 #coordinates of the image corners
12 pts_corners = np.array([[0,300,300,0],[0,0,200,200]])
13



```

```

14 #points on a square in the middle + image corners
15 pts_source = np.array([[50,150,150,50],[50,50,150,150]])
16 pts_source = np.concatenate((pts_source,pts_corners),axis=1)
17
18 #points on a diamond in the middle + image corners
19 pts_target = np.array([[100,160,100,40],[40,100,160,100]])
20 pts_target = np.concatenate((pts_target,pts_corners),axis=1)
21
22 #compute triangulation using mid-point between source and
23 #target to get triangles that are good for both.
24 pts_mid = 0.5*(pts_target+pts_source)
25 trimesh = Delaunay(pts_mid.transpose())
26 #we only need the vertex indices so extract them from
27 #the data structure returned by Delaunay
28 tri = trimesh.simplices.copy()
29
30 # display initial image
31 plt.imshow(image)
32 plt.triplot(pts_source[0,:],pts_source[1,:],tri,color='r',linewidth=2)
33 plt.plot(pts_source[0,:],pts_source[1:], 'ro')
34 plt.show()
35
36 # display warped image
37 (warped,tindex) = warp(image,pts_source,pts_target,tri)
38 plt.imshow(warped)
39 plt.triplot(pts_target[0,:],pts_target[1,:],tri,color='r',linewidth=2)
40 plt.plot(pts_target[0,:],pts_target[1:], 'ro')
41 plt.show()
42
43 # display animated movie by warping to weighted averages
44 # of pts_source and pts_target
45
46 #assemble an array of image frames
47 movie = []
48 for t in np.arange(0,1,0.1):
49     pts_warp = (1-t)*pts_source+t*pts_target
50     warped_image,tindex = warp(image,pts_source,pts_warp,tri)
51     movie.append(warped_image)
52
53 #use display_movie function defined in a5utils.py to create an animation
54 HTML(display_movie(movie), to_html())


```

```
54 from display_movie import CO_STREAM
55
56
```



Out[258]:





☐ Once ☒ Loop ☐ Reflect

<Figure size 432x288 with 0 Axes>

3. Face Morphing [15 pts]

Use your warping function in order to generate a morphing video between two faces. A separate notebook ***select_keypoints.ipynb*** has been provided that you can use to click keypoints on a pair of images in order to specify the correspondences. You should choose two color images of human faces to use (no animals or cartoons) and use the notebook interface to annotate corresponding keypoints on the two faces. To get a good result you should annotate 20-30 keypoints. The images should be centered on the faces with the face taking up most of the image frame. To keep the code simple, the two images should be the exact same dimension. Please use python or your favorite image editing tool to crop/scale them to the same size before you start annotating keypoints.

Once you have the keypoints saved, modify the code below to load in the keypoints and images, add the image corners to the set of points, and generate a morph sequence which starts with one face image and smoothly transitions to the other face image by simultaneously warping and cross-dissolving between the two.

To generate a frame of the morph at time t in the interval $[0,1]$, you should: (1) compute the intermediate shape as a weighted average of the keypoint locations of the two faces, (2) warp both image1 and image2 to this intermediate shape, (3) compute the weighted average of the two warped images.

You will likely want to refer to the code above for testing the ***warp*** function which is closely related.

For grading purposes, your notebook should display

1. The two images with keypoints and triangulations overlayed
2. Three intermediate frames of the morph sequence at $t=0.25$, $t=0.5$ and $t=0.75$

```
In [259]: 1 # load in the keypoints and images select_keypoints.ipynb
          2 f = open('face_correspondences.pckl', 'rb')
          3 image1, image2, pts1, pts2 = pickle.load(f)
          4
          5 # add the image corners as additional points so that the
          6 # triangles cover the whole image
          7 (h1, w1, c1) = image1.shape
          8 (h2, w2, c2) = image2.shape
          9
         10 corners1 = np.array([[0, w1, 0, w1], [0, 0, h1, h1]])
```

```


11 corners2 = np.array([[0, w2, 0, w2],[0, 0, h2, h2]])
12
13 pts1 = np.append(pts1, corners1, axis = 1)
14 pts2 = np.append(pts2, corners2, axis = 1)
15
16
17 #compute triangulation using mid-point between source and
18 #target to get trianglest that are good for both.
19 pts_mid = 0.5*(pts1+pts2)
20 trimesh = Delaunay(pts_mid.transpose())
21 tri = trimesh.simplices.copy()
22
23
24 # generate the frames of the morph
25 movie = []
26 for t in np.arange(0,1,0.05):
27     pts_warp = ((1-t)*pts1) + (t*pts2)
28     warped_image1,tindex1 = warp(image1, pts1, pts_warp, tri)
29     warped_image2,tindex2 = warp(image2, pts2, pts_warp, tri)
30     result_image = warped_image1*(1-t) + warped_image2*(t)
31     movie.append(result_image)
32
33
34
35
36 # display original images and overlaid triangulation
37 fig = plt.figure(figsize = (10,12))
38 ax1 = fig.add_subplot(1,2,1)
39 ax1.imshow(image1)
40 ax1.triplot(pts1[0,:],pts1[1,:],tri,color='g',linewidth=1)
41 ax1.plot(pts1[0,:],pts1[1:], 'r.')
42
43 ax2 = fig.add_subplot(1,2,2)
44 ax2.imshow(image2)
45 ax2.triplot(pts2[0,:],pts2[1,:],tri,color='g',linewidth=1)
46 ax2.plot(pts2[0,:],pts2[1:], 'r.')
47 plt.show()
48
49
50 # # display images at t=0.25, t=0.5 and t=0.75
51 # # i.e. visualize movie[5], movie[10], movie[15]

```

```

51 # i.e. visualize movie[5], movie[10], movie[15]
52 fig2 = plt.figure(figsize = (10,12))
53 fig2.add_subplot(1,3,1).imshow(movie[5])
54 fig2.add_subplot(1,3,2).imshow(movie[10])
55 fig2.add_subplot(1,3,3).imshow(movie[15])
56 plt.show()
57
58 # optional: display as an animated movie
59

```



In [262]:


```
1 HTML(display_movie(movie).to_jshtml())
```

Out[262]:



☒ Once ☐ Loop ☐ Reflect

<Figure size 432x288 with 0 Axes>



4. Face Swapping [15 pts]

We can use the same machinery of piecewise affine warping in order to swap faces. To accomplish this, we first annotate two faces with keypoints as we did for morphing. In this case they keypoints should only cover the face and we won't add the corners of the image. To place the face from image1 into image2, you should call your **warp** function to generate the warped face image1_warped. In order to composite only the warped face pixels, we need to create an alpha map. You can achieve this by using the **tindex** map returned from your warp function to make a binary mask which is True inside the face region and False else where. In order to

minimize visible artifacts, you should utilize **`scipy.ndimage.gaussian_filter`** in order to feather the edge of the alpha mask (as we did in a previous assignment for panorama mosaic blending). Once you have the feathered alpha map, you can composite the `image1_warped` face with the background from `image2`.

You should display in your submitted pdf notebook (1) the two source images with the keypoints overlaid, (2) the face from `image1` overlaid on `image2`, (3) the face from `image2` overlaid on `image1`.

It is *ok* to use the same faces for this part and the morphing part. However, to get the best results for face swapping it is important to only include keypoints inside the face while for morphing it may be better to include additional keypoints (e.g., in order to morph the hair, clothes etc.)

In [263]:

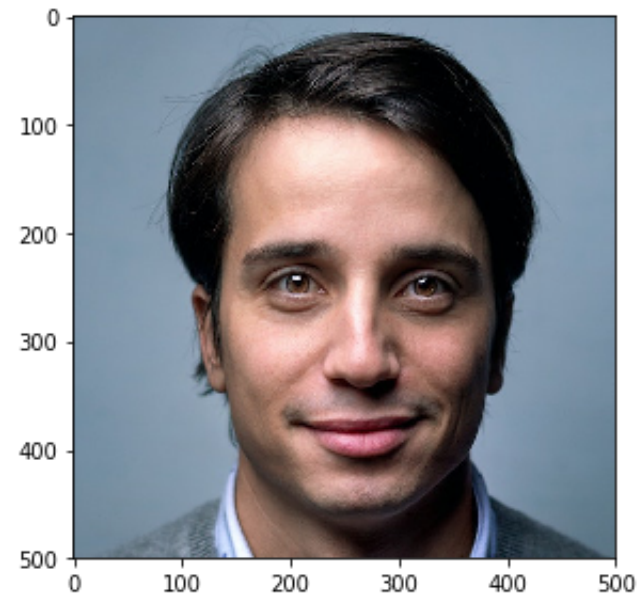
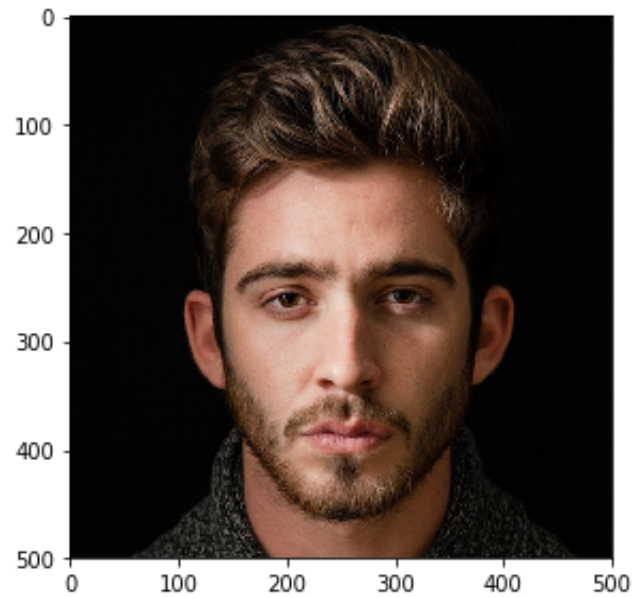
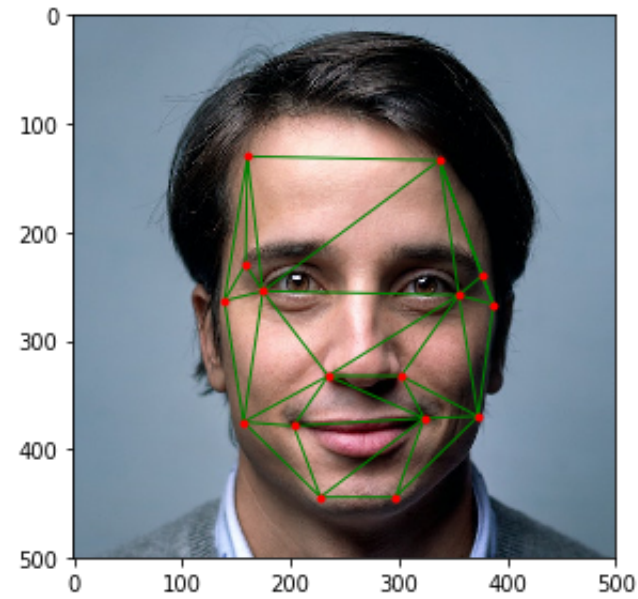
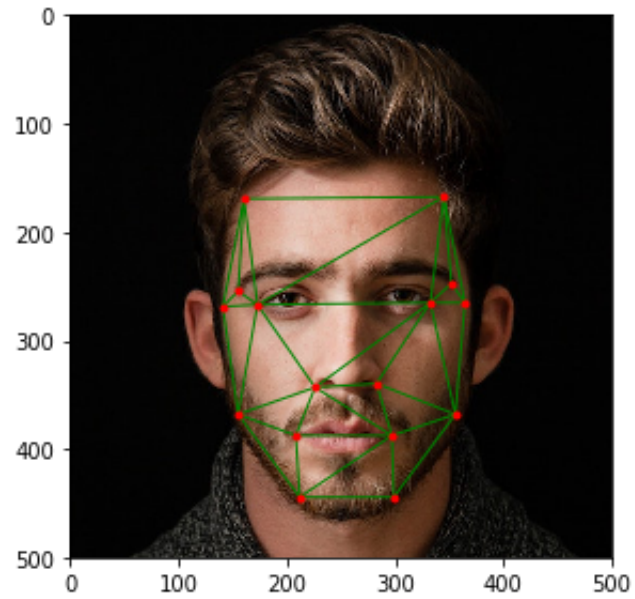
```
1  f = open('face_correspondences.pckl','rb')
2  image1,image2,pts1,pts2 = pickle.load(f)
3  f.close()
4
5  #compute triangulation using mid-point between source and
6  #target to get triangles that are good for both images.
7  pts_mid = 0.5*(pts1+pts2)
8  trimesh = Delaunay(pts_mid.transpose())
9  tri = trimesh.simplices.copy()
10
11 # put the face from image1 in to image2
12 (warped,tindex) = warp(image1, pts1, pts2, tri)
13
14 mask = np.where(tindex== -1, 0., 1)
15 alpha = gaussian_filter(mask, sigma = 8)
16 alpha = alpha*mask
17 alpha[np.where(alpha < 0)] = 0
18 alpha[np.where(alpha!=1)] -= np.min(alpha[np.nonzero(alpha)])
19 alpha[np.where(alpha<0)] = 0
20
21
22
23 swap1 = np.zeros(image1.shape)
24 # do an alpha blend of the warped image1 and image2
25 swap1[:, :, 0] = alpha * warped[:, :, 0] + (1.-alpha)*image2[:, :, 0]
26 swap1[:, :, 1] = alpha * warped[:, :, 1] + (1.-alpha)*image2[:, :, 1]
27 swap1[:, :, 2] = alpha * warped[:, :, 2] + (1.-alpha)*image2[:, :, 2]
28
```

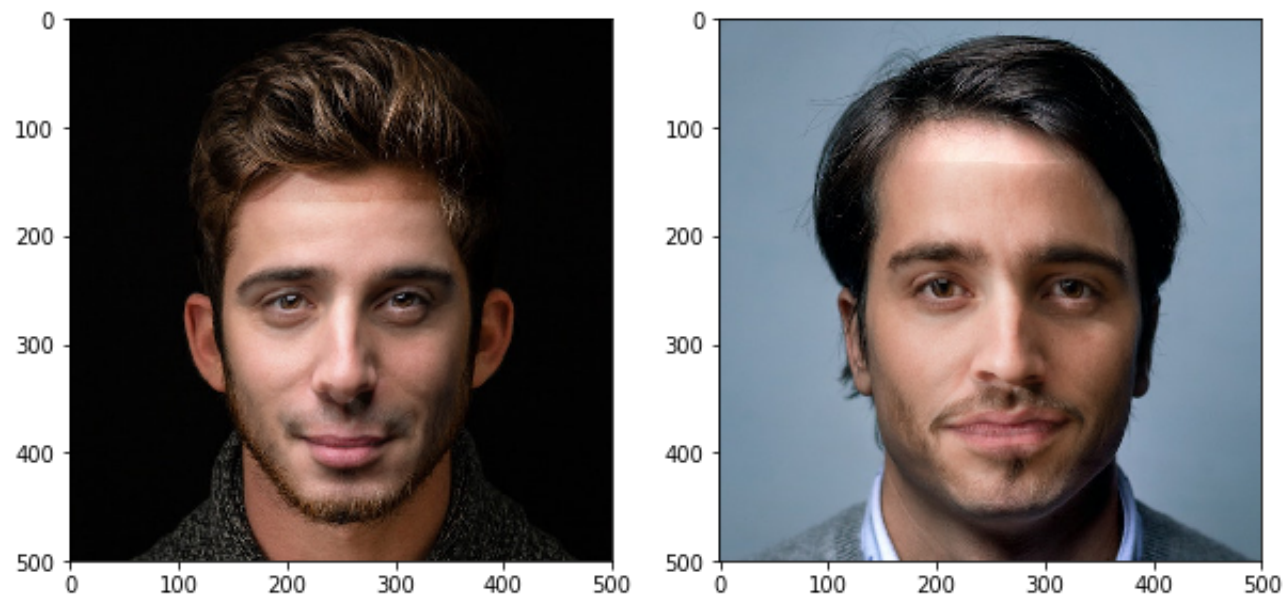
```

29
30 #now do the swap in the other direction
31
32 (warped,tindex) = warp(image2, pts2, pts1, tri)
33
34 mask = np.where(tindex==-1, 0., 1)
35 alpha = gaussian_filter(mask, sigma = 8)
36 alpha = alpha*mask
37 alpha[np.where(alpha < 0)] = 0
38 alpha[np.where(alpha!=1)] -= np.min(alpha[np.nonzero(alpha)])
39 alpha[np.where(alpha<0)] = 0
40
41 swap2 = np.zeros(image2.shape)
42 # do an alpha blend of the warped imagel and image2
43 swap2[:, :, 0] = alpha * warped[:, :, 0] + (1.-alpha)*image1[:, :, 0]
44 swap2[:, :, 1] = alpha * warped[:, :, 1] + (1.-alpha)*image1[:, :, 1]
45 swap2[:, :, 2] = alpha * warped[:, :, 2] + (1.-alpha)*image1[:, :, 2]
46
47
48 # alphaFig = plt.figure(figsize=(10,12))
49 # alphaFig.add_subplot(1,1,1).imshow(alpha)
50 # plt.show()
51
52
53 # display the images with the keypoints overlaid
54 fig = plt.figure(figsize = (10,12))
55 ax1 = fig.add_subplot(1,2,1)
56 ax1.imshow(image1)
57 ax1.triplot(pts1[0,:],pts1[1,:],tri,color='g',linewidth=1)
58 ax1.plot(pts1[0,:],pts1[1,:], 'r.')
59
60 ax2 = fig.add_subplot(1,2,2)
61 ax2.imshow(image2)
62 ax2.triplot(pts2[0,:],pts2[1,:],tri,color='g',linewidth=1)
63 ax2.plot(pts2[0,:],pts2[1,:], 'r.')
64
65
66 # display the face swapping result
67 fig3 = plt.figure(figsize = (10, 12))
68 fig3.add_subplot(2,2,1).imshow(image1)
69 fig3.add_subplot(2,2,2).imshow(image2)

```

```
70 fig3.add_subplot(2,2,3).imshow(swap2)
71 fig3.add_subplot(2,2,4).imshow(swap1)
72 plt.show()
73
```





In []:

1