



## Object Detection (due Saturday 3/9/2019)

In this assignment, you will develop an object detector based on gradient features and sliding window classification. A set of test images and *hogvis.py* are provided in the Canvas assignment directory

---

**Name:** Dikai Fang

**SID:** 29991751

---

In [102]:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
```

## 1. Image Gradients [20 pts]

Write a function that takes a grayscale image as input and returns two arrays the same size as the image, the first of which contains the magnitude of the image gradient at each pixel and the second containing the orientation.

Your function should filter the image with the simple x- and y-derivative filters described in class. Once you have the derivatives you can compute the orientation and magnitude of the gradient vector at each pixel. You should use ***scipy.ndimage.correlate*** with the 'nearest' option in order to nicely handle the image boundaries.



Include a visualization of the output of your gradient calculate for a small test image. For displaying the orientation result, please uses a cyclic colormap such as "hsv" or "twilight". (see <https://matplotlib.org/tutorials/colors/colormaps.html> (<https://matplotlib.org/tutorials/colors/colormaps.html>))

In [103]:

```
1  #we will only use:  scipy.ndimage.correlate
2  from scipy import ndimage
3
4  def mygradient(image):
5      """
6          This function takes a grayscale image and returns two arrays of the
7          same size, one containing the magnitude of the gradient, the second
8          containing the orientation of the gradient.
9
10
11      Parameters
12      -----
13      image : 2D float array of shape HxW
14              An array containing pixel brightness values
15
16      Returns
17      -----
18      mag : 2D float array of shape HxW
19              gradient magnitudes
20
21      ori : 2Dfloat array of shape HxW
22              gradient orientations in radians
23      """
24
25      # your code goes here
26      #      sobel_x = np.array([[1., 0 , -1],[2, 0, -2],[1, 0, -1]])
27      #      sobel_y = np.array([[1., 2, 1], [0, 0, 0], [-1, -2, -1]])
28
29      x_kernel = np.array([[-1, 1]])
30      y_kernel = np.array([[-1], [1]])
31
32      x_dir = ndimage.correlate(image, x_kernel, mode = "nearest") + 1e-16
33      y_dir = ndimage.correlate(image, y_kernel, mode = "nearest")
34
35      mag = ((x_dir**2) + (y_dir**2))**(1/2)
36      ori = np.arctan((y_dir/x_dir))
37      #      ori = np.arctan((y_dir/x_dir))
38
39      return (mag,ori)
```


In [104]:

```
1 #
2 # Demonstrate your mygradient function here by loading in a grayscale
3 # image, calling mygradient, and visualizing the resulting magnitude
4 # and orientation images. For visualizing orientation image, I suggest
5 # using the hsv or twilight colormap.
6 #
7
8 image = plt.imread("https://upload.wikimedia.org/wikipedia/commons/f/fa/Grayscale_8bits_palette
9
10 (mag,ori) = mygradient(image)
11
12 fig = plt.figure(figsize=(10,12))
13 fig.add_subplot(1,2,1).imshow(mag, cmap=plt.cm.gray)
14 fig.add_subplot(1,2,2).imshow(ori, cmap=plt.cm.gray)
15
16 plt.show()
17
18 #visualize results.
19
```



0 20 40 60 80 100 120 140

0 20 40 60 80 100 120 140



## 2. Histograms of Gradient Orientations [25 pts]

Write a function that computes gradient orientation histograms over each 8x8 block of pixels in an image. Your function should bin the orientation into 9 equal sized bins between  $-\pi/2$  and  $\pi/2$ . The input of your function will be an image of size  $H \times W$ . The output should be a three-dimensional array **ohist** whose size is  $(H/8) \times (W/8) \times 9$  where **ohist[i,j,k]** contains the count of how many edges of orientation  $k$  fell in block  $(i,j)$ . If the input image dimensions are not a multiple of 8, you should use **np.pad** with the **mode=edge** option to pad the width and height up to the nearest integer multiple of 8.

To determine if a pixel is an edge, we need to choose some threshold. I suggest using a threshold that is 10% of the maximum gradient magnitude in the image. Since each 8x8 block will contain a different number of edges, you should normalize the resulting histogram for each block to sum to 1 (i.e., **`np.sum(ohist,axis=2)`** should be 1 at every location).

I would suggest your function loops over the orientation bins. For each orientation bin you'll need to identify those pixels in the image whose magnitude is above the threshold and whose orientation falls in the given bin. You can do this easily in numpy using logical operations in order to generate an array the same size as the image that contains Trues at the locations of every edge pixel that falls in the given orientation bin and is above threshold. To collect up pixels in each 8x8 spatial block you can use the function **`ski.util.view_as_windows(...,(8,8),step=8)`** and **`np.count_nonzeros`** to count the number of edges in each block.

Test your code by creating a simple test image (e.g. a white disk on a black background), computing the descriptor and using the provided function **`hogvis`** to visualize it.

Note: in the discussion above I have assumed 8x8 block size and 9 orientations. In your code you should use the parameters **`bsize`** and **`norient`** in place of these constants.

In [105]:

```
1  #we will only use:  ski.util.view_as_windows for computing hog descriptor
2  import skimage as ski
3  #we will only use:  scipy.ndimage.correlate
4  from scipy import ndimage
5
6  def hog(image,bsize=8,norient=9):
7
8      """
9      This function takes a grayscale image and returns a 3D array
10     containing the histogram of gradient orientations descriptor (HOG)
11     We follow the convention that the histogram covers gradients starting
12     with the first bin at -pi/2 and the last bin ending at pi/2.
13
14     Parameters
15     -----
16     image : 2D float array of shape HxW
17           An array containing pixel brightness values
18
19     bsize : int
20           The size of the spatial bins in pixels, defaults to 8
21
22     norient : int
23           The number of orientation histogram bins defaults to 9
```



```

23     the number of orientation histogram bins, defaults to 5
24
25 Returns
26 -----
27 ohist : 3D float array of shape (H/bsize,W/bsize,norient)
28         edge orientation histogram
29
30 """
31
32 # determine the size of the HOG descriptor
33 (h,w) = image.shape
34 h2 = int(np.ceil(h/float(bsize)))
35 w2 = int(np.ceil(w/float(bsize)))
36 ohist = np.zeros((h2,w2,norient))
37
38 # pad the input image as needed so that it is a multiple of bsize
39 pw = 0
40 ph = 0
41 if (w % bsize != 0):
42     pw = bsize - (w % bsize)
43 if (h % bsize != 0):
44     ph = bsize - (h % bsize)
45 image = np.pad(image,((ph,0),(pw,0)),'edge')
46
47 # compute image gradients
48 (mag,ori) = mygradient(image)
49
50 # choose a threshold which is 10% of the maximum gradient magnitude in the image
51 thresh = np.amax(mag) * 0.1
52
53 # separate out pixels into orientation channels, dividing the range of orientations
54 # [-pi/2,pi/2] into norient equal sized bins and count how many fall in each block
55 # as a sanity check, make sure every pixel gets assigned to at most 1 bin.
56 bincount = np.zeros((h2*bsize,w2*bsize))
57
58 orient_start = (-np.pi/2)
59 orient_end = (np.pi/2)
60 orient_inter = (abs(orient_start) + abs(orient_end))/norient
61
62 for i in range(norient):
63     #create a binary image containing 1s for pixels at the ith

```

```

64         #orientation where the magnitude is above the threshold.
65     #         B = np.zeros((h, w))
66     B = np.zeros(image.shape)
67
68     lowerBound = orient_start + i * orient_inter
69     higherBound = orient_start + (i+1) * orient_inter
70
71     oneIdx = (ori >= lowerBound) & (ori < higherBound) & (mag > thresh)
72     if (i == norient-1):
73         oneIdx = (ori >= lowerBound) & (ori <= higherBound) & (mag > thresh)
74     B[oneIdx] = 1
75
76     #sanity check
77     bincount = bincount + B
78
79     #pull out non-overlapping bsize x bsize blocks
80     chblock = ski.util.view_as_windows(B, (bsize, bsize), step=bsize)
81     chblock = chblock.reshape((h2, w2, bsize*bsize))
82
83     #sum up the count for each block and store the results
84     ohist[:, :, i] = np.count_nonzero(chblock, axis = 2)
85
86
87     assert(np.all(bincount<=1))
88
89     # lastly, normalize the histogram so that the sum along the orientation dimension is 1
90     # note: don't divide by 0! If there are no edges in a block (i.e. the sum of counts
91     # is 0) then your code should leave all the values as zero.
92
93     total = np.sum(ohist, axis = 2)
94     total[np.where(total == 0)] = 1.
95
96     for i in range(norient):
97         ohist[:, :, i] = ohist[:, :, i]/total
98
99     assert(ohist.shape==(h2,w2,norient))
100
101     return ohist

```






In [ ]:

1

In [106]:

```
1  #provided function for visualizing hog descriptors
2  import hogvis as hogvis
3
4  #
5  # generate a simple test image... a 80x80 image
6  # with a circle of radius 30 in the center
7  #
8  [yy,xx] = np.mgrid[-40:40,-40:40]
9  # im = np.array((xx*xx+yy*yy<=30*30),dtype=float)
10
11 # [yy,xx] = np.mgrid[-44:44,-44:44]
12 im = np.array((xx*xx+yy*yy<=30*30),dtype=float)
13
14 hog_info = hog(im)
15 hogvis_result = hogvis.hogvis(hog_info, bsize=8, norient=9)
16
17
18 #
19 # display the image and the output of hogvis
20 #
21
22 fig = plt.figure(figsize = (10,12))
23
24 fig.add_subplot(1,2,1).imshow(im)
25 fig.add_subplot(1,2,2).imshow(hogvis_result)
26
27 plt.show()
28
29
```





### 3. Detection [25 pts]

Write a function that takes a template and an image and returns the top detections found in the image. Your function should follow the definition given below.

In your function you should first compute the histogram-of-gradient-orientation feature map for the image, then correlate the template with the feature map. Since the feature map and template are both three dimensional, you will want to filter each orientation separately and then sum up the results to get the final response. If the image of size  $H \times W$  then this final response map will be of size  $(H/8) \times (W/8)$ .

When constructing the list of top detections, your code should implement non-maxima suppression so that it doesn't return overlapping detections. You can do this by sorting the responses in descending order of their score. Every time you add a detection to the list to return, check to make sure that the location of this detection is not too close to any of the detections already in the output list. You can estimate the overlap by computing the distance between a pair of detections and checking that the distance is greater than say 70% of the width of the template.

Your code should return the locations of the detections in terms of the original image pixel coordinates (so if your detector had a high response at block  $[i,j]$  in the response map, then you should return  $(8i,8j)$  as the pixel coordinates).

I have provided a function for visualizing the resulting detections which you can use to test your detect function. Please include some visualization of a simple test case.

```
In [107]: 1 def detect(image, template, ndetect=5, bsize=8, norient=9):
          2     """
          3
```

```

3
4     This function takes a grayscale image and a HOG template and
5     returns a list of detections where each detection consists
6     of a tuple containing the coordinates and score (x,y,score)
7
8     Parameters
9     -----
10    image : 2D float array of shape HxW
11            An array containing pixel brightness values
12
13    template : a 3D float array
14            The HOG template we wish to match to the image
15
16    ndetect : int
17            Number of detections to return
18
19    bsize : int
20            The size of the spatial bins in pixels, defaults to 8
21
22    norient : int
23            The number of orientation histogram bins, defaults to 9
24
25    Returns
26    -----
27    detections : a list of tuples of length ndetect
28                Each detection is a tuple (x,y,score)
29
30    """
31
32    # norient for the template should match the norient parameter passed in
33    assert(template.shape[2]==norient)
34
35    fmap = hog(image,bsize=bsize,norient=norient)
36
37
38    #cross-correlate the template with the feature map to get the total response
39    # resp = np.zeros(fmap.shape)
40    resp = np.zeros((fmap.shape[0], fmap.shape[1]))
41    for i in range(norient):
42        resp = resp + ndimage.correlate(fmap[:, :, :, i], template[:, :, :, i], mode = "nearest")
43

```

```

44     #sort the values in resp in descending order.
45     # val[i] should be ith largest score in resp
46     # ind[i] should be the index at which it occurred so that val[i]==resp[ind[i]]
47     #
48     #     val = ... #sorted response values
49     #     ind = ... #corresponding indices
50
51     val = np.sort(resp, axis = None)[::-1]
52     ind = np.unravel_index(np.argsort(resp, axis=None)[::-1], resp.shape)
53
54
55     #work down the list of responses from high to low, to generate a
56     # list of ndetect top scoring matches which do not overlap
57     detcount = 0
58     i = 0
59     detections = []
60     while ((detcount < ndetect) and (i < len(val))):
61         # convert 1d index into 2d index
62         yb = ind[0][i]
63         xb = ind[1][i]
64
65         assert(val[i]==resp[yb,xb]) #make sure we did indexing correctly
66
67         #covert block index to pixel coordinates based on bsize
68         xp = bsize * xb
69         yp = bsize * yb
70
71         #check if this detection overlaps any detections that we've already added
72         #to the list. compare the x,y coordinates of this detection to the x,y
73         #coordinates of the detections already in the list and see if any overlap
74         #by checking if the distance between them is less than 70% of the template
75         # width/height
76
77         overlap = False
78         for d in range(len(detections)):
79             dist = abs((yp - detections[d][1])**2 + (xp - detections[d][0])**2)**(1/2)
80             x_dist = abs((xp - detections[d][0]))
81             y_dist = abs((yp - detections[d][1]))
82             #         if ((dist < 0.7 * template.shape[0]) or (dist < 0.7 * template.shape[1])):
83                 if ((x_dist < 0.7 * template.shape[1]) or (y_dist < 0.7 * template.shape[0])):
84                     overlap = True

```

```

85         break
86
87
88         #if the detection doesn't overlap then add it to the list
89         if (not overlap):
90             detcount = detcount + 1
91             detections.append((xp,yp,val[i]))
92
93         i=i+1
94
95     if (len(detections) < ndetect):
96         print('WARNING: unable to find ',ndetect,' non-overlapping detections')
97
98     return detections

```

In [ ]:

1

In [108]:

```

1  import matplotlib.patches as patches
2
3  def plot_detections(image,detections,tsize_pix):
4      """
5      This is a utility function for visualization that takes an image and
6      a list of detections and plots the detections overlayed on the image
7      as boxes.
8
9      Color of the bounding box is based on the order of the detection in
10     the list, fading from green to red.
11
12     Parameters
13     -----
14     image : 2D float array of shape HxW
15             An array containing pixel brightness values
16
17     detections : a list of tuples of length ndetect
18                  Detections are tuples (x,y,score)
19
20     tsize_pix : (int,int)
21                 The height and width of the box in pixels
22
23     Returns

```

```

23         return
24         -----
25     None
26
27     """
28     ndetections = len(detections)
29
30     fig = plt.figure()
31
32     #     plt.imshow(image, aspect = 'equal')
33     fig.add_subplot(1,1,1).imshow(image)
34     ax = plt.gca()
35     w = tsize_pix[1]
36     h = tsize_pix[0]
37     red = np.array([1,0,0])
38     green = np.array([0,1,0])
39     ct = 0
40     for (x,y,score) in detections:
41         xc = x-(w//2)
42         yc = y-(h//2)
43         col = (ct/ndetections)*red + (1-(ct/ndetections))*green
44         rect = patches.Rectangle((xc,yc),w,h,linewidth=3,edgecolor=col,facecolor='none')
45         ax.add_patch(rect)
46         ct = ct + 1
47
48     plt.show()

```

In [109]:

```


1  #
2  # sketch of some simple test code, modify as needed
3  #
4
5  #create a synthetic image
6  [yy,xx] = np.mgrid[-40:40,-40:40]
7  im1 = np.array((xx*xx+yy*yy<=30*30),dtype=float)
8  [yy,xx] = np.mgrid[-60:20,-60:20]
9  im2 = np.array((xx*xx+yy*yy<=25*25),dtype=float)
10 im = 0.5*im1+0.5*im2
11
12 #compute feature map with default parameters
13 fmap = hog(im)
14

```

```

15 #extract a 3x3 template
16 # tplate = fmap[2:5,2:5,:]
17 template = fmap[1:4,1:4,:]
18
19 #run the detect code
20 detections = detect(im,template,ndetect=5)
21
22 #visualize results.
23 plot_detections(im,detections,(24,24))
24
25 # visually confirm that:
26 # 1. top detection should be the same as the location where we selected the template
27 # 2. multiple detections do not overlap too much

```



## 4. Learning Templates [15 pts]

The final step is to implement a function to learn a template from positive and negative examples. Your code should take a collection of cropped positive and negative examples of the object you are interested in detecting, extract the features for each, and generate a template by taking the average positive template minus the average negative template.

```

1 def learn_template(positive_examples, negative_examples, num_pos=10, num_neg=10, num_iter=100):

```



```

in [110]: 1 def learn_template(posfiles,negfiles,tsize=np.array([10,10]),bsize=8,norient=9):
2         """
3         This function takes a list of positive images that contain cropped
4         examples of an object + negative files containing cropped background
5         and a template size. It produces a HOG template and generates visualization
6         of the examples and template
7
8         Parameters
9         -----
10        posfiles : list of str
11                   Image files containing cropped positive examples
12
13        negfiles : list of str
14                   Image files containing cropped negative examples
15
16        tsize : (int,int)
17                The height and width of the template in blocks
18
19        Returns
20        -----
21        template : float array of size tsize x norient
22                   The learned HOG template
23
24        """
25
26        #compute the template size in pixels
27        #corresponding to the specified template size (given in blocks)
28        tsize_pix=bsize*tsize
29
30        #figure to show positive training examples
31        fig1 = plt.figure()
32        pltct = 1
33
34        #accumulate average positive and negative templates
35        pos_t = np.zeros((tsize[0],tsize[1],norient),dtype=float)
36        for file in posfiles:
37            #load in a cropped positive example
38            img = plt.imread(file)
39
40            #convert to grayscale and resize to fixed dimension tsize_pix
41            grayI = np.zeros((img.shape[0], img.shape[1]), dtype = float)

```

```

42     grayI[:, :, :] = (0.299 * img[:, :, 0] + 0.587 * img[:, :, 1] + 0.114 * img[:, :, 2],
43
44     #using skimage.transform.resize if needed.
45     grayI = ski.transform.resize(grayI, tsize)
46
47     #display the example if you want to train with a large # of examples,
48     #you may want to modify this, e.g. to show only the first 5.
49     ax = fig1.add_subplot(len(posfiles), 1, pltct)
50     ax.imshow(grayI, cmap=plt.cm.gray)
51     pltct = pltct + 1
52
53     #extract feature
54     fmap = hog(grayI, bsize=1)
55
56     #compute running average
57     # pos_t = ..
58
59     pos_t += fmap.reshape(pos_t.shape)
60
61     pos_t = (1/len(posfiles))*pos_t
62     fig1.show()
63
64
65
66     # repeat same process for negative examples
67     fig2 = plt.figure()
68     pltct = 1
69     neg_t = np.zeros((tsize[0], tsize[1], norient), dtype=float)
70     for file in negfiles:
71         img = plt.imread(file)
72
73         grayI = np.zeros((img.shape[0], img.shape[1]), dtype = np.float64)
74         grayI[:, :, :] = (0.299 * img[:, :, 0] + 0.587 * img[:, :, 1] + 0.114 * img[:, :, 2],
75         grayI = ski.transform.resize(grayI, tsize)
76
77         ax = fig2.add_subplot(len(negfiles), 1, pltct)
78         ax.imshow(grayI, cmap=plt.cm.gray)
79         pltct = pltct + 1
80
81         fmap = hog(grayI, bsize = 1)
82

```

```
83         neg_t += fmap
84
85
86     neg_t = (1/len(negfiles))*neg_t
87     fig2.show()
88
89     # add code here to visualize the positive and negative parts of the template
90     # using hogvis. you should separately visualize pos_t and neg_t rather than
91     # the final template.
92     hogFig = plt.figure()
93     pos_hogvis_result = hogvis.hogvis(pos_t, bsize=8, norient=9)
94     neg_hogvis_result = hogvis.hogvis(neg_t, bsize=8, norient=9)
95     pos_ax = hogFig.add_subplot(1,2,1)
96     pos_ax.imshow(pos_hogvis_result)
97     pos_ax.set_title("positive template")
98     neg_ax = hogFig.add_subplot(1,2,2)
99     neg_ax.imshow(neg_hogvis_result)
100    neg_ax.set_title("negative template")
101    hogFig.show()
102
103    # now construct our template as the average positive minus average negative
104    template = pos_t - neg_t
105
106    return template
107
```

## 5. Experiments [15 pts]

Test your detection by training a template and running it on a test image.

In your experiments and writeup below you should include: (a) a visualization of the positive and negative patches you use to train the template and corresponding hog feature, (b) the detection results on the test image. You should show (a) and (b) for **two different object categories**, the provided face test images and another category of your choosing (e.g. feel free to experiment with detecting cat faces, hands, cups, chairs or some other type of object). Additionally, please include results of testing your detector where there are at least 3 objects to detect (this could be either 3 test images which each have one or more objects, or a single image with many (more than 3) objects). Your test image(s) should be distinct from your training examples. Finally, write a brief (1 paragraph) discussion of where the detector works well and when it fails. Describe some ways you might be able to make it better.

NOTE 1: You will need to create the cropped test examples to pass to your **learn\_template**. You can do this by cropping out the examples by hand (e.g. using an image editing tool). You should attempt to crop them out in the most consistent way possible, making sure that each example is centered with the same size and aspect ratio. Negative examples can be image patches that don't contain the object of interest. You should crop out negative examples with roughly the same resolution as the positive examples.

NOTE 2: For the best result, you will want to test on images where the object is the same size as your template. I recommend using the default **bsize** and **norient** parameters for all your experiments. You will likely want to modify the template size as needed

### Experiment 1: Face detection



In [111]:


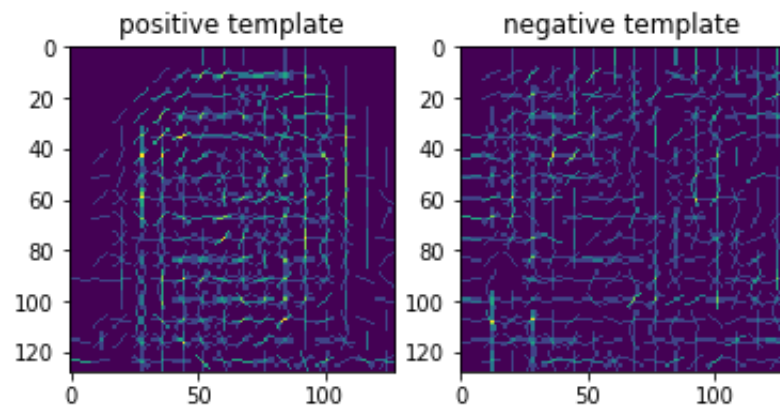

```
1  # assume template is 16x16 blocks, you may want to adjust this
2  # for objects of different size or aspect ratio.
3  # compute image a template size
4  bsize=8
5  tsize=np.array([16,16]) #height and width in blocks
6  tsize_pix = bsize*tsize #height and width in pixels
7  posfiles = ('pos1.jpg', 'pos2.jpg', 'pos3.jpg', 'pos4.jpg', 'pos5.jpg')
8  negfiles = ('neg1.jpg', 'neg2.jpg', 'neg3.jpg', 'neg4.jpg', 'neg5.jpg')
9
10 # call learn_template to learn and visualize the template and training data
```


```

10 # call learn_template to learn and visualize the template and training data
11 template = learn_template(posfiles,negfiles,tsize=tsize)
12
13 # call detect on one or more test images, visualizing the result with the plot_detections funct
14
15 im = plt.imread('images/faces/faces3.jpg')
16 grayIm = np.zeros((im.shape[0], im.shape[1]), dtype = float)
17 grayIm[:, :] = (0.299 * im[:, :, 0] + 0.587 * im[:, :, 1] + 0.114 * im[:, :, 2])
18 detections = detect(grayIm, template, ndetect = 4)
19 plot_detections(im,detections,tsize_pix)
20
21
22 im2 = plt.imread('images/faces/faces5.jpg')
23 grayIm2 = np.zeros((im2.shape[0], im2.shape[1]), dtype = float)
24 grayIm2[:, :] = (0.299 * im2[:, :, 0] + 0.587 * im2[:, :, 1] + 0.114 * im2[:, :, 2])
25 detections2 = detect(grayIm2, template, ndetect = 7)
26 plot_detections(im2,detections2,tsize_pix)
27

```







## Experiment 2: ??? detection

In [112]:

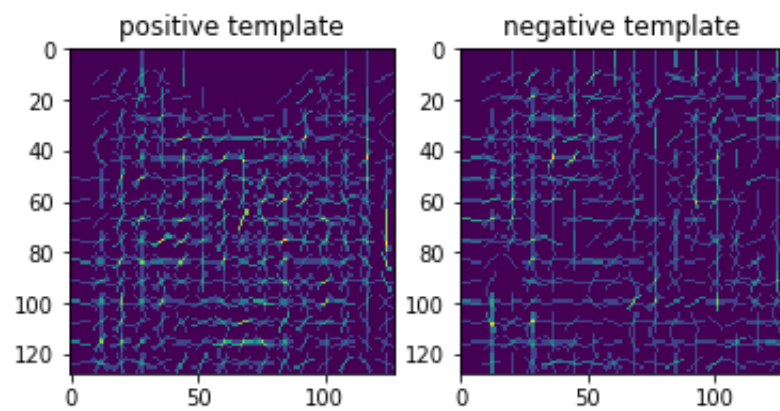


```


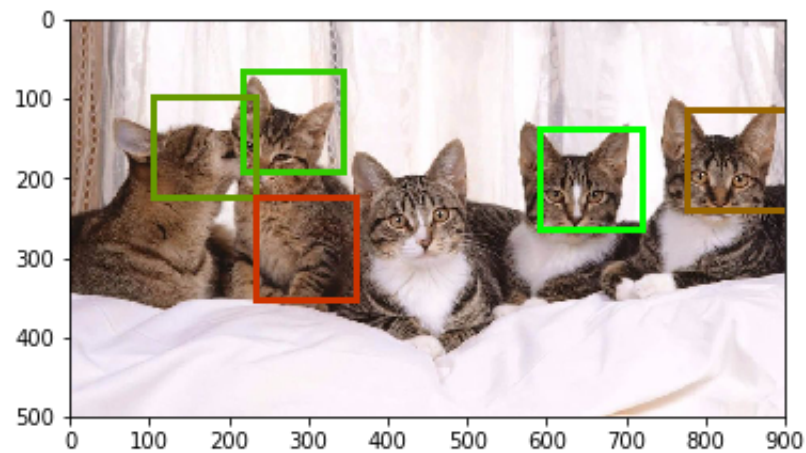
1  bsize=8
2  tsize=np.array([16,16]) #height and width in blocks
3  tsize_pix = bsize*tsize #height and width in pixels
4  posfiles = ('catPos1.jpg', 'catPos2.jpg', 'catPos3.jpg', 'catPos4.jpg', 'catPos5.jpg')
5  negfiles = ('neg1.jpg', 'neg2.jpg', 'neg3.jpg', 'neg4.jpg', 'neg5.jpg')
6
7
8  template = learn_template(posfiles,negfiles,tsize=tsize)
9
10
11 im = plt.imread('images/cats/cats1.jpg')
12 grayIm = np.zeros((im.shape[0], im.shape[1]), dtype = float)
13 grayIm[:, :] = (0.299 * im[:, :, 0] + 0.587 * im[:, :, 1] + 0.114 * im[:, :, 2])
14 detections = detect(grayIm, template)
15 plot_detections(im,detections,tsize_pix)
16
17
18 im2 = plt.imread('images/cats/cats2.jpg')
19 grayIm2 = np.zeros((im2.shape[0], im2.shape[1]), dtype = float)
20 grayIm2[:, :] = (0.299 * im2[:, :, 0] + 0.587 * im2[:, :, 1] + 0.114 * im2[:, :, 2])
21 detections2 = detect(grayIm2, template)
22 plot_detections(im2,detections2,tsize_pix)
23
24

```









The final result is depending on many facts, such as the training images and block size. Training images and the test image are more similar and consistent, the more faces can be detected. For example, face4.jpg has upside down faces and face2.jpg has varied facial expression. These faces can barely be detected because I only used regular faces. Block size also needs to adjust when the test image's resolution changes. When the image has higher resolution, the block size becomes smaller. Thus detection fails since it cannot find the target.